



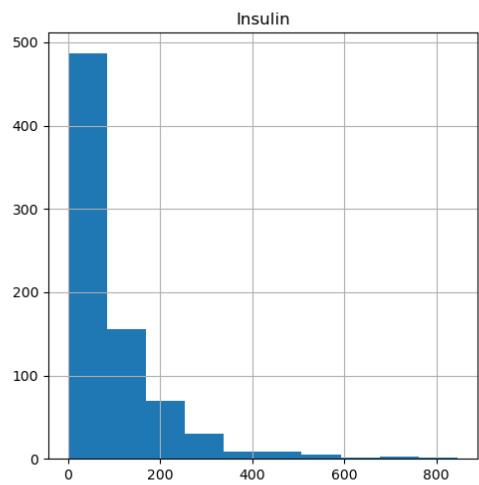






# Data Exploration

```
In [1]: 1 import pandas as pd
        2 diabetes_data = pd.read_csv('diabetes.csv')
        3 p = diabetes_data.hist(figsize = (20,20))
        4 diabetes_data.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```



In [2]:

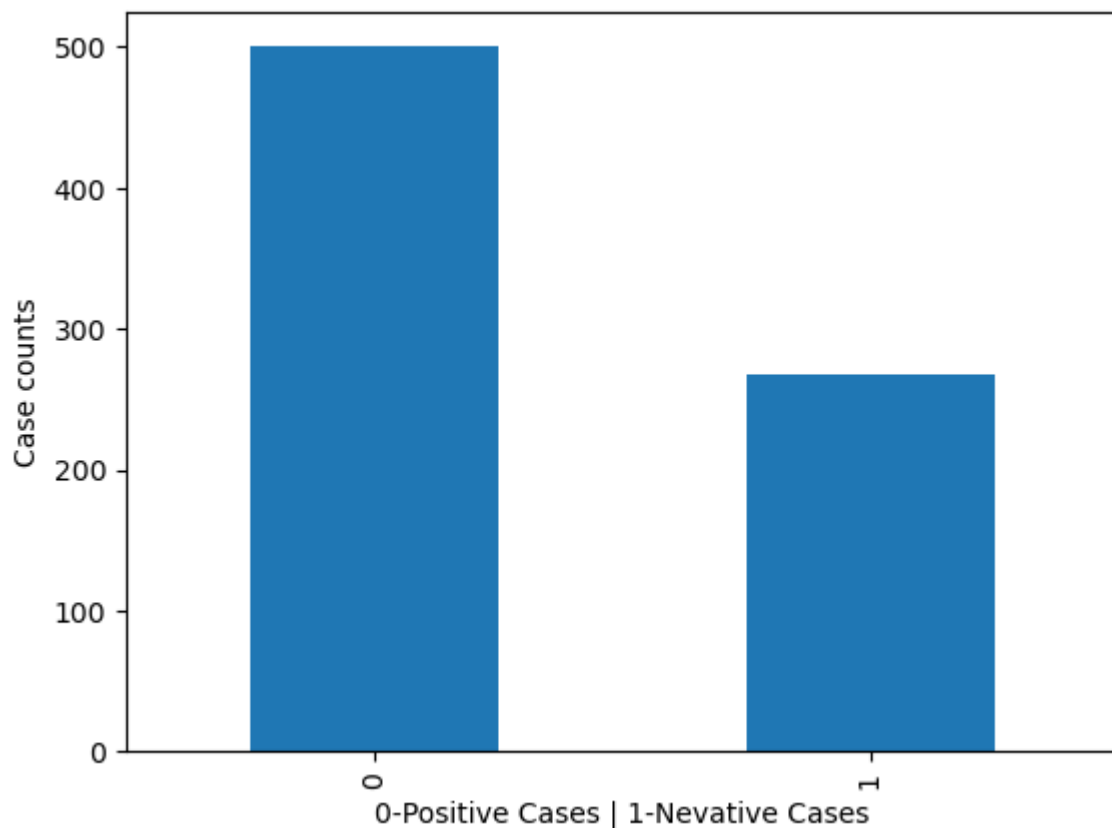
```
1
2  ## checking the balance of the data by plotting the count of outcomes by the
3  color_wheel = {1: "#0392cf",
4                 2: "#7bc043"}
5  colors = diabetes_data["Outcome"].map(lambda x: color_wheel.get(x + 1))
6  print(diabetes_data.Outcome.value_counts())
7  p=diabetes_data.Outcome.value_counts().plot(kind="bar",)
8  p.set_xlabel("0-Positive Cases | 1-Negative Cases")
9  p.set_ylabel("Case counts")
```

```
0    500
```

```
1    268
```

```
Name: Outcome, dtype: int64
```

```
Text(0, 0.5, 'Case counts')
```



```
In [3]: 1 from sklearn.datasets import load_svmlight_file
2 diabetes_features, diabetes_labels = load_svmlight_file('diabetes_scale.txt')
3 #Since the data is of the type SVMLight file using a tht particular reader to
4 print(diabetes_features.shape)
```

(768, 8)

```
In [4]: 1 #the data in of the form of a tuple
2 diabetes_features = diabetes_features.toarray()
3 print(type(diabetes_features), type(diabetes_labels))
```

<class 'numpy.ndarray'> <class 'numpy.ndarray'>

```
In [5]: 1 import numpy as np
2 import random
3 np.random.seed(10)
4 weights = np.random.rand(8)
5 # weights = np.array([0.,0.,0.,0.,0.,0.,0.,0.])
6 weights
```


array([0.77132064, 0.02075195, 0.63364823, 0.74880388, 0.49850701,  
0.22479665, 0.19806286, 0.76053071])

In [6]:

```
1 def accuracy(y_true, y_pred):
2     accuracy = np.mean(y_pred == y_true)
3     return accuracy
4 #Predicting for each row {-1,1} and then getting tht into pred array and re
5 def predict(features,weight):
6     pred = list()
7     for row in features:
8         activation = 0
9         for i in range(len(row)):
10             activation += weights[i ] * row[i] #calc Xi*Wi
11             pred.append(np.sign(activation))# Predictig for 1 row
12     return (accuracy(diabetes_labels, pred))
13 def update_weights(weights):
14     l_rate = 0.0001
15     summation = 0
16     for row,y in zip(diabetes_features,diabetes_labels):
17         # calaculating the summation of 1*8 vector for each row
18         summation+= (y*row*(1 if y*(np.dot(row,weights))<0 else 0))
19     #now modifying the weights according to learning for
20     weights += (l_rate*summation)
21     # print("updated",weights)
22     return weights
```

In [7]:

```
1  # print("1st accuracy" ,predict(diabetes_features,weights))
2  import matplotlib.pyplot as plt
3  best_weights = 0
4  best_accu = 0
5  acc_list,loss_list = [],[]
6  for i in range(40):
7      acc = predict(diabetes_features,weights)
8      weights = update_weights(weights)
9      acc_list.append(acc)
10     if acc>best_accu:
11         best_accu = acc
12         best_weights = weights
13     # l = zero_one_loss(diabetes_features,diabetes_labels,weights)
14     # loss_list.append(l)
15     # plt.figure(figsize=(5, 4))
16     # plt.title("Plotting Loss with 20 Epochs ")
17     # plt.plot(np.arange(20), e_train_loss_list, label = " Loss")
18     # plt.legend()
19     # plt.show()
20
21     #Plotting Validation Curve
22     plt.figure(figsize=(5, 4))
23     plt.title("Plotting Accuracy with 40 Epochs")
24     plt.plot(np.arange(40), acc_list, label = "Accuracy")
25     # plt.plot(np.arange(40), loss_list, label = "Loss")
26     plt.legend()
27     plt.show()
28
29     print("best accuracy: ",best_accu)
30     print("best weights: " ,best_weights)
```



best accuracy: 0.55078125

best weights: [ 0.10378526 -0.10073597 0.51230559 0.24198056 -0.09956228 0.03309548  
-0.2585228 0.05204732]

## Experimentation

In [8]:

```
1 #Adding an additional column to the dataset initialized as all 1 & and addit
2 import numpy as np
3 type(diabetes_features)
4 bias_col = np.ones(len(diabetes_features)).reshape(-1,1)
5 mod_diabetes_features = np.concatenate([diabetes_features, bias_col], axis=1)
6 print(mod_diabetes_features.shape)
```

(768, 9)



```

In [9]: import train_test_split
        1
        2
        3 = train_test_split(mod_diabetes_features, diabetes_labels, test_size=0.2,
        4 train_test_split(X_bigtrain, y_bigtrain, test_size=0.20, random_state=42,
        5
        6
        7
        8
        9
        10
        11 train_test_split(diabetes_features, diabetes_labels, test_size=0.20, random_state=42,
        12 stratify=diabetes_labels)
        13
        14
        15
        16

```

Splits With Bias

Training Size : (491, 9)

Testing Size : (154, 9)

Validation Size : (123, 9)

Splits With Bias

Training Size : (491, 8)

Testing Size : (154, 8)

Validation Size : (123, 8)

In [10]:

```
1 import seaborn as sns
2 def accuracy(y_true, y_pred):
3     accuracy = np.mean(y_pred == y_true)
4     return accuracy
5 #Predicting for each row {-1,1} and then getting tht into pred array and re
6 def predict(features,labels,weights):
7     pred = list()
8     for row in features:
9         activation = 0
10        for i in range(len(row)):
11            activation += weights[i] * row[i] #calc Xi*Wi
12        pred.append(np.sign(activation))# Predictig for 1 row
13    return (accuracy(labels, pred))
14 def accuracy_scores(features,labels,weights):
15     from sklearn.metrics import accuracy_score, f1_score, precision_score, r
16     from sklearn.metrics import roc_curve,roc_auc_score
17     pred = list()
18     for row in features:
19         activation = 0
20         for i in range(len(row)):
21             activation += weights[i] * row[i] #calc Xi*Wi
22         pred.append(np.sign(activation))# Predictig for 1 row
23     print("F1 Score :",f1_score(labels, pred, average="binary"))
24     print("Precision :",precision_score(labels, pred, average="binary"))
25     print("Recall :",recall_score(labels, pred, average="binary"))
26     print("Confusion Matrix :")
27     cmat = confusion_matrix(labels, pred)
28     sns.heatmap(cmat,annot=True)
29     ax= plt.subplot()
30     sns.heatmap(cmat, annot=True, fmt='g', ax=ax);
31     # labels, title and ticks
32     plt.title('Confusion Matrix')
33     ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
34     ax.set_title('Confusion Matrix');
35     ax.xaxis.set_ticklabels(['Positive', 'Negative']); ax.yaxis.set_ticklabe
36
37 #     plt.set_xticklabels(['Positive', 'Negative'])
38 #     plt.set_yticklabels(['Positive', 'Negative'])
```

```

39     plt.show()
40     return
41 def update_weights(weights,features,labels,loss,l_rate):
42     #     l_rate = 0.001
43     summation = 0
44     for row,y in zip(features,labels):
45         # calaculating the summation of 1*8 vector for each row
46         if(loss == "perceptron"):
47             ls = max(0,-1*y*(np.dot(row,weights))) #Perceptron Loss
48             summation+= (y*row*ls)
49         elif(loss == "zero_one"):
50             ls = (1 if y*(np.dot(row,weights))<0 else 0) #ZeroOne Loss
51             summation+= (y*row*ls)
52         #now modifying the weights according to learning for
53         weights += (l_rate*summation)
54     #     print("updated",weights)
55     return weights
56 def zero_one_loss(features,labels,weights):
57     loss =0
58     for row ,y in zip(features,labels):
59         loss+=(1 if y*(np.dot(row,weights))<0 else 0)
60     return loss/len(labels)
61
62 def perceptron_loss(features,labels,weights):
63     loss =0
64     for row ,y in zip(features,labels):
65         loss+= max(0,-1*y*(np.dot(row,weights)))
66     return loss/len(labels)

```

## SK learn baseline

```
In [11]: 1 from sklearn.linear_model import Perceptron
2 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
3 clf = Perceptron(tol=1e-3, random_state=42)
4 clf.fit(X_bigtrain, y_bigtrain)
5 Perceptron()
6 sk_pred = clf.predict(X_test)
7 print("Testing Accuracy for Sklearn Baseline : ",clf.score(X_test, y_test))
8 print("F1 Score :",f1_score(y_test, sk_pred, average="binary"))
9 print("Precision :",precision_score(y_test, sk_pred, average="binary"))
10 print("Recall :",recall_score(y_test, sk_pred, average="binary"))
```

```
Testing Accuracy for Sklearn Baseline : 0.7987012987012987
F1 Score : 0.8472906403940887
Precision : 0.8349514563106796
Recall : 0.86
```

```
In [12]: 1 #Initializing 8 weights + 1 Bias
2 np.random.seed(89)
3 b_weights = np.random.rand(9)
4 # weights = np.zeros(8)
5 weights
```

```
array([ 0.10378526, -0.10073597,  0.51230559,  0.24198056, -0.09956228,
        0.03309548, -0.2585228 ,  0.05204732])
```

## Experimenting with initial weights with bias

```

In [13]: 1 #Experimenting with the the best Random seed for weights
2 all_acc=list()
3 overall_best_acc = 0
4 overall_best_weights = 0
5 l_rate =0.001
6 for j in range(202):
7     np.random.seed(j)
8     b_weights = np.random.rand(9)
9     best_weights = 0
10    best_accu = 0
11    for i in range(20):
12        b_weights = update_weights(b_weights,b_X_train,b_y_train,"zero_one",l_rate)
13        train_acc = predict(b_X_train,b_y_train,b_weights)
14        val_acc = predict(b_X_Val,b_y_Val,b_weights)
15        #         print("Iteration ",i)
16        #         print("Training accuracy:",train_acc," Validation accuracy:",val_acc)
17        if val_acc>best_accu:
18            best_accu = val_acc
19            best_weights = b_weights
20        #     print("for Random Seed:",j)
21        #     print("best accuracy: ",best_accu)
22        if best_accu>overall_best_acc:
23            overall_best_acc = best_accu
24            overall_best_weights = best_weights
25        all_acc.append((best_accu,j))
26        #     print("best weights: " ,best_weights)
27    all_acc.sort(reverse=True)
28    all_acc[:6]

```

```

[(0.7886178861788617, 89),
 (0.7560975609756098, 111),
 (0.7479674796747967, 195),
 (0.7398373983739838, 10),
 (0.7317073170731707, 160),
 (0.7317073170731707, 107)]

```

## Experimenting with initial weights without bias

```

In [14]: 1 #Experimenting with the the best Random seed for weights
2 all_acc=list()
3 overall_best_acc = 0
4 overall_best_weights = 0
5 l_rate =0.001
6 for j in range(202):
7     np.random.seed(j)
8     weights = np.random.rand(8)
9     best_weights = 0
10    best_accu = 0
11    for i in range(20):
12        b_weights = update_weights(weights,X_train,y_train,"zero_one",l_rate)
13        train_acc = predict(X_train,y_train,weights)
14        val_acc = predict(X_Val,y_Val,weights)
15        #         print("Iteration ",i)
16        #         print("Training accuracy:",train_acc," Validation accuracy:",val_acc)
17        if val_acc>best_accu:
18            best_accu = val_acc
19            best_weights = weights
20        #     print("for Random Seed:",j)
21        #     print("best accuracy: ",best_accu)
22        if best_accu>overall_best_acc:
23            overall_best_acc = best_accu
24            overall_best_weights = best_weights
25        all_acc.append((best_accu,j))
26        #     print("best weights: " ,best_weights)
27    all_acc.sort(reverse=True)
28    all_acc[:6]

```

```

[(0.7723577235772358, 84),
 (0.7723577235772358, 33),
 (0.7723577235772358, 26),
 (0.7642276422764228, 195),
 (0.7642276422764228, 140),
 (0.7642276422764228, 138)]

```

## Hyper-Parameter Tuning with Bias

In [15]:

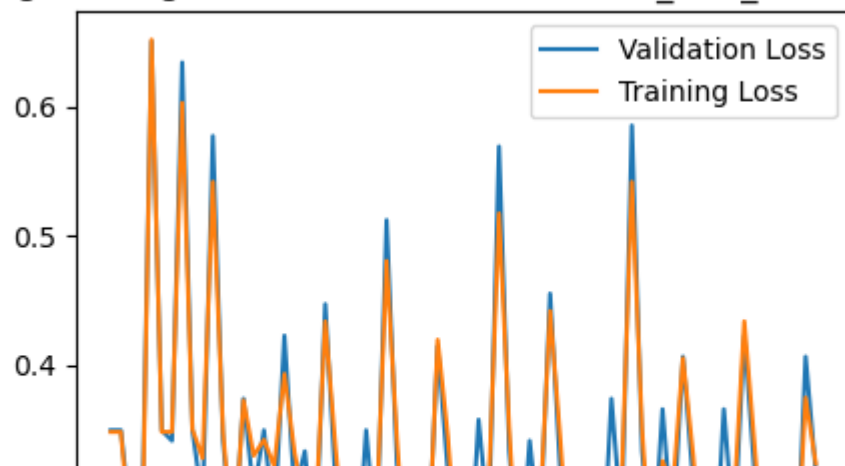
```
1 import matplotlib.pyplot as plt
2 b_accuracy_table = list()
3 learning_Rates = [1,.1,.01,.001,0.0001]
4 epoch = np.arange(10,101,10)
5 loss =["perceptron","zero_one"]
6 for e in epoch:
7     for lr_rate in learning_Rates:
8         for l in loss:
9             best_weights = 0
10            best_epoch = 0
11            best_accu = 0
12            e_val_loss_list = []
13            e_train_loss_list = []
14            e_val_acc_list = []
15            e_train_acc_list = []
16            np.random.seed(89)
17            b_weights = np.random.rand(9)
18            for i in range(e):
19                b_weights = update_weights(b_weights,b_X_train,b_y_train,l,lr_rate)
20                train_acc = predict(b_X_train,b_y_train,b_weights)
21                val_acc = predict(b_X_Val,b_y_Val,b_weights)
22                e_val_acc_list.append(val_acc)
23                e_train_acc_list.append(train_acc)
24                if l == "zero_one":
25                    e_train_loss = zero_one_loss(b_X_train,b_y_train,b_weights)
26                    e_val_loss = zero_one_loss(b_X_Val,b_y_Val,b_weights)
27                else:
28                    e_train_loss = perceptron_loss(b_X_train,b_y_train,b_weights)
29                    e_val_loss = perceptron_loss(b_X_Val,b_y_Val,b_weights)
30                e_val_loss_list.append(e_val_loss)
31                e_train_loss_list.append(e_train_loss)
32            # print("Iteration ",i)
33            # print("Training accuracy:",train_acc," Validation accuracy:",val_acc)
34            if val_acc>best_accu:
35                best_epoch = (i+1)
36                best_accu = val_acc
37                best_weights = b_weights
38
```

```


39     b_accuracy_table.append((best_accu,lr_rate,l,best_epoch,e,best_w
40     #Plotting Loss Curve
41     plt.figure(figsize=(5, 4))
42     if l == "zero_one":
43         plt.title("Plotting Training vs Validation Loss with Zero_One
44     else:
45         plt.title("Plotting Training vs Validation Loss with Percepti
46     plt.plot(np.arange(e), e_val_loss_list, label = "Validation Loss'
47     plt.plot(np.arange(e), e_train_loss_list, label = "Training Loss'
48     plt.legend()
49     plt.show()
50
51     #Plotting Validation Curve
52     plt.figure(figsize=(5, 4))
53     if l == "zero_one":
54         plt.title("Plotting Training vs Validation Accuracy with Zero
55     else:
56         plt.title("Plotting Training vs Validation Accuracy with Perc
57
58     plt.plot(np.arange(e), e_val_acc_list, label = "Validation Accur
59     plt.plot(np.arange(e), e_train_acc_list, label = "Training Accur
60     plt.legend()
61     plt.show()
62     print("best accuracy: ",best_accu)
63     print("best weights: " ,best_weights)

```

Plotting Training vs Validation Loss with Zero\_One\_Loss with lr = 1







In [16]:

```
1 import pandas as pd
2 b_accuracy_table.sort(reverse=True)
3 dfb = pd.DataFrame(b_accuracy_table, columns=["Val_Accuracy", "Learning Rate"
4 dfb.iloc[:45,:]
```

	Val_Accuracy	Learning Rate	Loss	nth Epoch	Epoch Range	
0	0.788618	1.000	zero_one	84	100	[-130.2063526795285, -43
1	0.788618	1.000	zero_one	84	90	[-202.20635517952866, -3
2	0.788618	0.100	zero_one	76	100	[-13.647380979530904, -4
3	0.788618	0.100	zero_one	76	90	[-9.370908279530937, -37
4	0.788618	0.100	zero_one	76	80	[-10.200318579530958, -3
5	0.788618	0.010	zero_one	66	100	[-0.8885427295311398, -4
6	0.788618	0.010	zero_one	66	90	[-0.9320719795311407, -4
7	0.788618	0.010	zero_one	66	80	[-1.1103071395311424, -4
8	0.788618	0.010	zero_one	66	70	[-1.9785424695311433, -3
9	0.788618	0.001	zero_one	20	100	[-0.18030591503116294, -
10	0.788618	0.001	zero_one	20	90	[-0.13636471053116314, -
11	0.788618	0.001	zero_one	20	80	[-0.14807057653116332, -
12	0.788618	0.001	zero_one	20	70	[-0.17307056103116342, -
13	0.788618	0.001	zero_one	20	60	[-0.13748229953116364, -
14	0.788618	0.001	zero_one	20	50	[-0.16436464153116373, -
15	0.788618	0.001	zero_one	20	40	[-0.16795285853116382, -
16	0.788618	0.001	zero_one	20	30	[-0.12395283553116401, -
17	0.788618	0.001	zero_one	20	20	[-0.05665868603116414, -
18	0.780488	1.000	zero_one	68	80	[-99.02984967952895, -36
19	0.780488	1.000	zero_one	68	70	[-171.20632117952903, -3
20	0.772358	1.000	zero_one	30	60	[-139.6180611795291, -43
21	0.772358	1.000	zero_one	30	50	[-95.7945091795294, -384
22	0.772358	1.000	zero_one	30	40	[-215.26510517952954, -3
23	0.772358	1.000	zero_one	30	30	[-128.73565967952982, -4
24	0.772358	0.100	zero_one	37	70	[-13.288552929530972, -3
25	0.772358	0.100	zero_one	37	60	[-19.794435179530986, -3
26	0.772358	0.100	zero_one	37	50	[-9.235607979530997, -43
27	0.772358	0.100	zero_one	37	40	[-7.653253329531025, -36
28	0.772358	0.010	zero_one	36	60	[-1.4632480595311448, -4
29	0.772358	0.010	zero_one	36	50	[-0.8938360345311473, -4
30	0.772358	0.010	zero_one	36	40	[-2.233836134531148, -3.5
31	0.764228	0.100	zero_one	10	30	[-13.800311929531034, -3

	Val_Accuracy	Learning Rate	Loss	nth Epoch	Epoch Range	
32	0.764228	0.100	zero_one	10	20	[-8.25913242953105, -40.7]
33	0.764228	0.100	zero_one	10	10	[-13.812073079531075, -2]
34	0.756098	1.000	zero_one	4	20	[-106.44152017953003, -4]
35	0.756098	1.000	zero_one	4	10	[-140.67680517953025, -2]
36	0.756098	0.010	zero_one	9	30	[-0.9632474545311513, -4]
37	0.756098	0.010	zero_one	9	20	[-0.8191295845311528, -4]
38	0.756098	0.010	zero_one	9	10	[-0.08089404453115911, -]
39	0.682927	0.001	zero_one	9	10	[0.08345899046883516, -]
40	0.650407	1.000	perceptron	1	100	[1.4688639284294427e+2]
41	0.650407	1.000	perceptron	1	90	[3.7459274320849006e+2]
42	0.650407	1.000	perceptron	1	80	[9.552942280671046e+22]
43	0.650407	1.000	perceptron	1	70	[2.436211268701485e+19]
44	0.650407	1.000	perceptron	1	60	[6.212876798970035e+17]

```
In [17]: 1 overall_best_weights = dfb['Best Weights'][0]
2 test_acc = predict(b_X_test,b_y_test,overall_best_weights)
3 print("Testing accuracy:",test_acc)
4 accuracy_scores(b_X_test,b_y_test,overall_best_weights)
```


Testing accuracy: 0.7727272727272727

F1 Score : 0.8275862068965517

Precision : 0.8155339805825242

Recall : 0.84

Confusion Matrix :



## Hyperparameter Tuning Without Bias


In [18]:

```
1 accuracy_table = list()
2 learning_Rates = [1,.1,.01,.001,0.0001]
3 epoch = np.arange(10,101,10)
4 loss = ["perceptron","zero_one"]
5 for e in epoch:
6     for lr_rate in learning_Rates:
7         for l in loss:
8             best_weights = 0
9             best_accu = 0
10            best_epoch = 0
11            e_val_loss_list = []
12            e_train_loss_list = []
13            e_val_acc_list = []
14            e_train_acc_list = []
15            np.random.seed(84)
16            weights = np.random.rand(8)
17            for i in range(e):
18                weights = update_weights(weights,X_train,y_train,l,lr_rate)
19                train_acc = predict(X_train,y_train,weights)
20                val_acc = predict(X_Val,y_Val,weights)
21                e_val_acc_list.append(val_acc)
22                e_train_acc_list.append(train_acc)
23                if l == "zero_one":
24                    e_train_loss = zero_one_loss(X_train,y_train,weights)
25                    e_val_loss = zero_one_loss(X_Val,y_Val,weights)
26                else:
27                    e_train_loss = perceptron_loss(X_train,y_train,weights)
28                    e_val_loss = perceptron_loss(X_Val,y_Val,weights)
29                e_val_loss_list.append(e_val_loss)
30                e_train_loss_list.append(e_train_loss)
31                # print("Iteration ",i)
32                # print("Training accuracy:",train_acc," Validation accuracy.
33                if val_acc>best_accu:
34                    best_epoch=(i+1)
35                    best_accu = val_acc
36                    best_weights = weights
37
38            accuracy_table.append((best_accu,lr_rate,l,best_epoch,e,best_weights))
```

```

39     #Plotting Loss Curve
40     plt.figure(figsize=(5, 4))
41     if l == "zero_one":
42         plt.title("Plotting Training vs Validation Loss with Zero_One")
43     else:
44         plt.title("Plotting Training vs Validation Loss with Perceptron_Loss")
45     plt.plot(np.arange(e), e_val_loss_list, label = "Validation Loss")
46     plt.plot(np.arange(e), e_train_loss_list, label = "Training Loss")
47     plt.legend()
48     plt.show()
49
50     #Plotting Validation Curve
51     plt.figure(figsize=(5, 4))
52     if l == "zero_one":
53         plt.title("Plotting Training vs Validation Accuracy with Zero_One")
54     else:
55         plt.title("Plotting Training vs Validation Accuracy with Perceptron_Loss")
56
57     plt.plot(np.arange(e), e_val_acc_list, label = "Validation Accuracy")
58     plt.plot(np.arange(e), e_train_acc_list, label = "Training Accuracy")
59     plt.legend()
60     plt.show()
61     print("best accuracy: ", best_accu)
62     print("best weights: " , best_weights)


```



best accuracy: 0.6910569105691057

best weights: [-0.164712 0.07884335 0.1884986 0.06869608 0.29625329 -0.08263069  
-0.23240394 -0.06289264]

Plotting Training vs Validation Loss with Perceptron\_Loss with  $l_r = 1$



```
In [19]: 1 import pandas as pd
2 accuracy_table.sort(reverse=True)
3 df = pd.DataFrame(accuracy_table, columns=["Val_Accuracy", "Learning Rate","I
4 df
```

	Val_Accuracy	Learning Rate	Loss	nth Epoch	Epoch Range	
0	0.796748	0.0010	zero_one	24	100	[-0.05242974486851241, -
1	0.796748	0.0010	zero_one	24	90	[-0.14925328036851254, -
2	0.796748	0.0010	zero_one	24	80	[-0.17001797686851272, -
3	0.796748	0.0010	zero_one	24	70	[-0.1546650178685129, -0
4	0.796748	0.0010	zero_one	24	60	[-0.07995910486851306, -
...	...	...	...	...	...	...
95	0.479675	0.0001	zero_one	20	20	[-0.2744531867685141, 0.
96	0.479675	0.0001	perceptron	30	30	[-0.25348978923140203, (
97	0.341463	0.0001	perceptron	20	20	[-0.23006416383874562, (
98	0.333333	0.0001	zero_one	1	10	[-0.13851198281851426, (
99	0.333333	0.0001	perceptron	1	10	[-0.15418712858391273, (

100 rows x 6 columns

```
In [20]: 1 overall_best_weights = df['Best Weights'][0]
2 test_acc = predict(X_test,y_test,overall_best_weights)
3 print("Testing accuracy:",test_acc)
4 accuracy_scores(X_test,y_test,overall_best_weights)
```

Testing accuracy: 0.564935064935065

F1 Score : 0.5109489051094891

Precision : 0.9459459459459459

Recall : 0.35

Confusion Matrix :

